

titanic-assignment

July 27, 2024

```
[1]: import pandas as pd
      from sklearn.model_selection import train_test_split
      from sklearn.preprocessing import StandardScaler, LabelEncoder
      import numpy as np
      import warnings
      warnings.filterwarnings("ignore")
      import matplotlib.pyplot as plt
      %matplotlib inline
      # magic command
      import seaborn as sns
```

```
[2]: import os
```

```
[3]: os.getcwd()
```

```
[3]: 'C:\\Users\\abhil'
```

```
[4]: df = pd.read_csv('titanic_train.csv')
```

```
[5]: df.head()
```

```
[5]: PassengerId  Survived  Pclass  \
0             1         0        3
1             2         1        1
2             3         1        3
3             4         1        1
4             5         0        3
```

```

                                Name      Sex  Age  SibSp  \
0                Braund, Mr. Owen Harris   male  22.0     1
1  Cumings, Mrs. John Bradley (Florence Briggs Th... female  38.0     1
2                Heikkinen, Miss. Laina   female  26.0     0
3  Futrelle, Mrs. Jacques Heath (Lily May Peel)   female  35.0     1
4                Allen, Mr. William Henry   male  35.0     0
```

```

   Parch      Ticket    Fare Cabin Embarked
0      0   A/5 21171   7.2500   NaN        S
```

1	0	PC	17599	71.2833	C85	C
2	0	STON/O2.	3101282	7.9250	NaN	S
3	0		113803	53.1000	C123	S
4	0		373450	8.0500	NaN	S

```
[6]: df.tail()
```

```
[6]:
```

	PassengerId	Survived	Pclass	Name \
886	887	0	2	Montvila, Rev. Juozas
887	888	1	1	Graham, Miss. Margaret Edith
888	889	0	3	Johnston, Miss. Catherine Helen "Carrie"
889	890	1	1	Behr, Mr. Karl Howell
890	891	0	3	Dooley, Mr. Patrick

	Sex	Age	SibSp	Parch	Ticket	Fare	Cabin	Embarked
886	male	27.0	0	0	211536	13.00	NaN	S
887	female	19.0	0	0	112053	30.00	B42	S
888	female	NaN	1	2	W./C. 6607	23.45	NaN	S
889	male	26.0	0	0	111369	30.00	C148	C
890	male	32.0	0	0	370376	7.75	NaN	Q

```
[7]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null    int64
1   Survived        891 non-null    int64
2   Pclass          891 non-null    int64
3   Name            891 non-null    object
4   Sex             891 non-null    object
5   Age            714 non-null    float64
6   SibSp          891 non-null    int64
7   Parch          891 non-null    int64
8   Ticket          891 non-null    object
9   Fare           891 non-null    float64
10  Cabin          204 non-null    object
11  Embarked       889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[8]: df.describe()
```

```
[8]:
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000

mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[9]: df.describe(include="O")# object
```

	Name	Sex	Ticket	Cabin	Embarked
count	891	891	891	204	889
unique	891	2	681	147	3
top	Braund, Mr. Owen Harris	male	347082	B96 B98	S
freq	1	577	7	4	644

```
[10]: df.describe(percentiles=[0.10,0.20,0.50,0.90])
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
10%	90.000000	0.000000	1.000000	14.000000	0.000000
20%	179.000000	0.000000	1.000000	19.000000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
90%	802.000000	1.000000	3.000000	50.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
10%	0.000000	7.550000
20%	0.000000	7.854200
50%	0.000000	14.454200

```

90%      2.000000   77.958300
max      6.000000  512.329200

```

```
[11]: df.isnull()
```

```

[11]:      PassengerId  Survived  Pclass   Name    Sex    Age  SibSp  Parch  Ticket  \
0             False     False   False  False  False  False  False  False  False
1             False     False   False  False  False  False  False  False  False
2             False     False   False  False  False  False  False  False  False
3             False     False   False  False  False  False  False  False  False
4             False     False   False  False  False  False  False  False  False
..            ...      ...    ...    ...    ...    ...    ...    ...
886          False     False   False  False  False  False  False  False  False
887          False     False   False  False  False  False  False  False  False
888          False     False   False  False  False   True  False  False  False
889          False     False   False  False  False  False  False  False  False
890          False     False   False  False  False  False  False  False  False

      Fare  Cabin  Embarked
0     False   True     False
1     False  False     False
2     False   True     False
3     False  False     False
4     False   True     False
..      ...    ...      ...
886  False   True     False
887  False  False     False
888  False   True     False
889  False  False     False
890  False   True     False

[891 rows x 12 columns]

```

```
[12]: df.isnull().sum()
```

```

[12]: PassengerId      0
Survived              0
Pclass                0
Name                  0
Sex                   0
Age                  177
SibSp                 0
Parch                 0
Ticket                0
Fare                  0
Cabin                687
Embarked              2

```

dtype: int64

```
[13]: df.isnull().sum().sum()
```

```
[13]: 866
```

```
[14]: df.isnull().any()
```

```
[14]: PassengerId    False
      Survived     False
      Pclass       False
      Name         False
      Sex          False
      Age          True
      SibSp        False
      Parch        False
      Ticket       False
      Fare         False
      Cabin        True
      Embarked     True
      dtype: bool
```

```
[15]: len(df)
```

```
[15]: 891
```

```
[16]: df1=df.copy()
```

```
[17]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Name            891 non-null   object
4   Sex             891 non-null   object
5   Age             714 non-null   float64
6   SibSp           891 non-null   int64
7   Parch           891 non-null   int64
8   Ticket          891 non-null   object
9   Fare            891 non-null   float64
10  Cabin           204 non-null   object
11  Embarked        889 non-null   object
```

```
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[18]: # percentage of missing value
df.isnull().sum()/len(df)*100
```

```
[18]: PassengerId      0.000000
Survived           0.000000
Pclass            0.000000
Name              0.000000
Sex               0.000000
Age              19.865320
SibSp            0.000000
Parch            0.000000
Ticket           0.000000
Fare             0.000000
Cabin            77.104377
Embarked         0.224467
dtype: float64
```

```
[19]: df1=df.copy()
```

```
[20]: df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null   int64
1   Survived         891 non-null   int64
2   Pclass           891 non-null   int64
3   Name             891 non-null   object
4   Sex              891 non-null   object
5   Age              714 non-null   float64
6   SibSp            891 non-null   int64
7   Parch            891 non-null   int64
8   Ticket           891 non-null   object
9   Fare             891 non-null   float64
10  Cabin            204 non-null   object
11  Embarked         889 non-null   object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[21]: df1.drop(['Name', 'Ticket', 'Cabin'], axis=1, inplace=True)
```

```
[22]: df1["Age"]=df1["Age"].fillna(df1["Age"].mean())
```

```
[23]: df1["Embarked"] = df1["Embarked"].fillna(df1["Embarked"].mode()[0])
```

```
[24]: df1.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId     891 non-null   int64
1   Survived        891 non-null   int64
2   Pclass          891 non-null   int64
3   Sex             891 non-null   object
4   Age             891 non-null   float64
5   SibSp           891 non-null   int64
6   Parch           891 non-null   int64
7   Fare            891 non-null   float64
8   Embarked        891 non-null   object
dtypes: float64(2), int64(5), object(2)
memory usage: 62.8+ KB
```

```
[25]: df1.isnull().sum().sum()
```

```
[25]: 0
```

```
[26]: plt.figure(figsize=(7,7))
sns.boxplot(data=df1)
```

```
[26]: <Axes: >
```



```
[27]: df1.describe()
```

```
[27]:
```

	PassengerId	Survived	Pclass	Age	SibSp	\
count	891.000000	891.000000	891.000000	891.000000	891.000000	
mean	446.000000	0.383838	2.308642	29.699118	0.523008	
std	257.353842	0.486592	0.836071	13.002015	1.102743	
min	1.000000	0.000000	1.000000	0.420000	0.000000	
25%	223.500000	0.000000	2.000000	22.000000	0.000000	
50%	446.000000	0.000000	3.000000	29.699118	0.000000	
75%	668.500000	1.000000	3.000000	35.000000	1.000000	
max	891.000000	1.000000	3.000000	80.000000	8.000000	

	Parch	Fare
count	891.000000	891.000000


mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[28]: df1['Sex'] = df1['Sex'].map({'male': 0, 'female': 1})
df1['Embarked'] = df1['Embarked'].map({'C': 0, 'Q': 1, 'S': 2})
```

```
[29]: def cap_outliers(series, lower_percentile=0.05, upper_percentile=0.95):
    lower_bound = series.quantile(lower_percentile)
    upper_bound = series.quantile(upper_percentile)
    return np.clip(series, lower_bound, upper_bound)
numerical_columns = ['Age', 'SibSp', 'Parch', 'Fare']
for column in numerical_columns:
    df1[column] = cap_outliers(df1[column])
```

```
[30]: plt.figure(figsize=(9,9))
sns.boxplot(data=df1)
```

```
[30]: <Axes: >
```






```
[31]: df1.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PassengerId  891 non-null    int64
1   Survived    891 non-null    int64
2   Pclass      891 non-null    int64
3   Sex         891 non-null    int64
4   Age         891 non-null    float64
5   SibSp       891 non-null    int64
6   Parch       891 non-null    int64
```

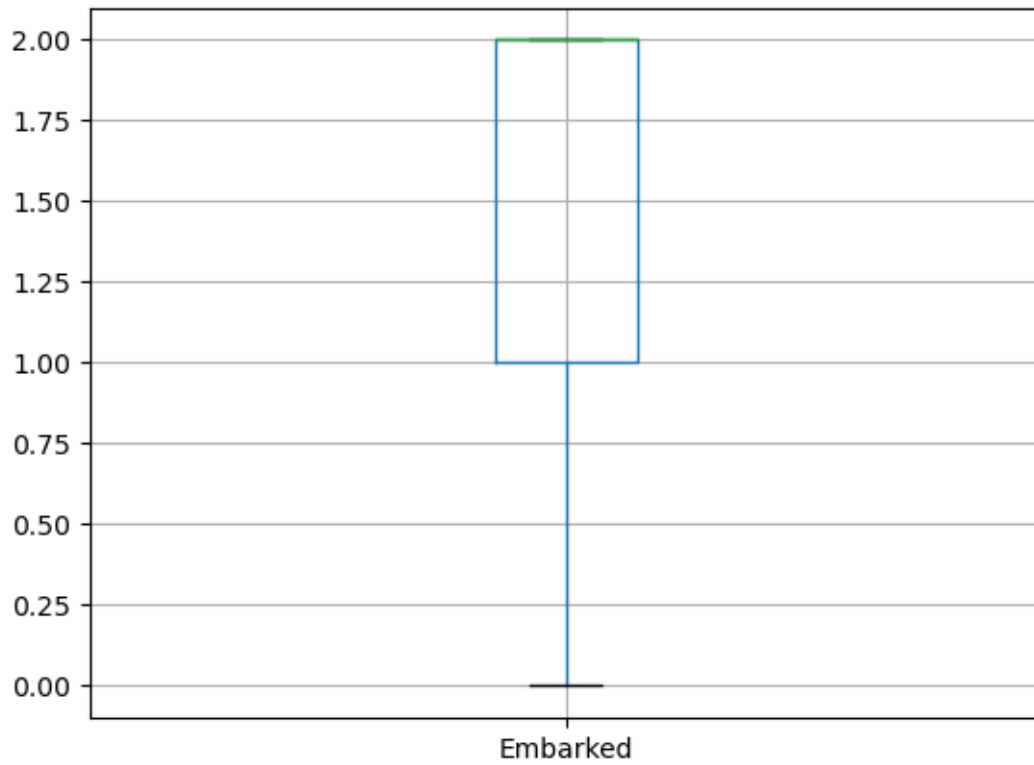
```
7   Fare      891 non-null   float64
8   Embarked  891 non-null   int64
dtypes: float64(2), int64(7)
memory usage: 62.8 KB
```

```
[32]: df100=df1[["Age","SibSp","Parch","Fare", "Embarked"]]
      for column in df100:
          plt.figure()
          df100.boxplot([column])
```












```
[33]: plt.figure(figsize=(10, 6))
sns.histplot(df1['Age'], kde=True)
plt.title('Age Distribution')
plt.show()
```




```
[34]: plt.figure(figsize=(10, 6))
sns.histplot(df1['Fare'], kde=True)
plt.title('Fare Distribution')
plt.show()
```



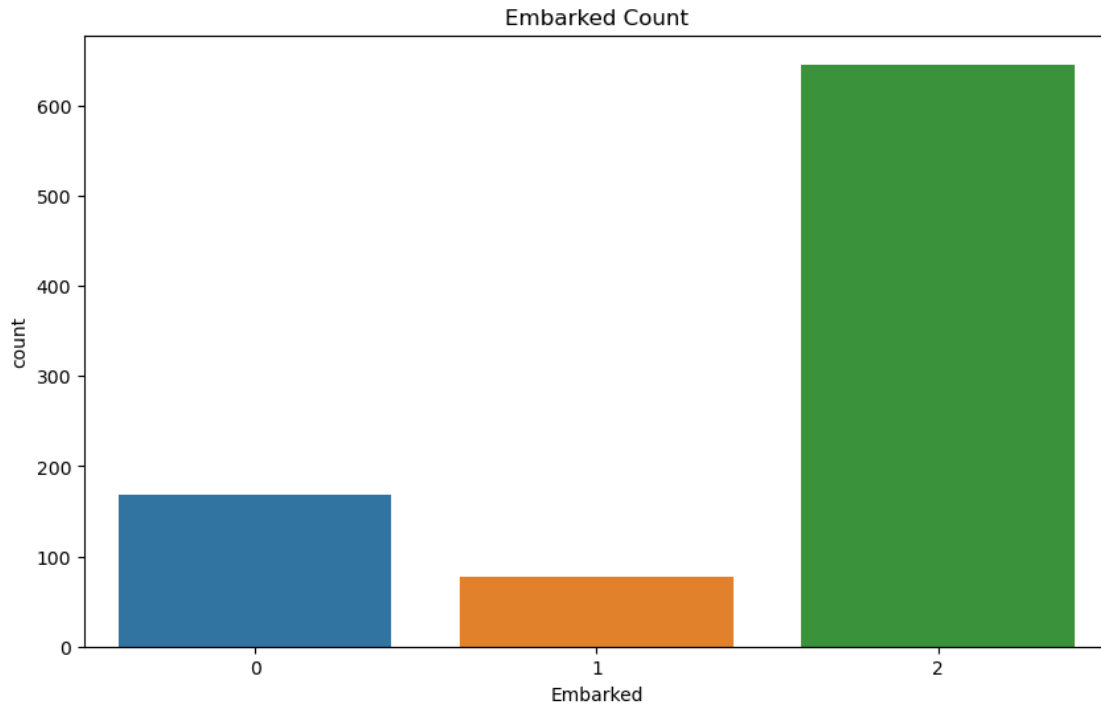
```
[35]: plt.figure(figsize=(10, 6))
sns.countplot(x='Survived', data=df1)
plt.title('Survival Count')
plt.show()
```

```
[36]: plt.figure(figsize=(10, 6))
sns.countplot(x='Pclass', data=df1)
plt.title('Passenger Class Count')
plt.show()
```



```
[37]: plt.figure(figsize=(10, 6))
sns.countplot(x='Embarked', data=df1)
plt.title('Embarked Count')
plt.show()
```



```
[38]: features = ['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare', 'Embarked']

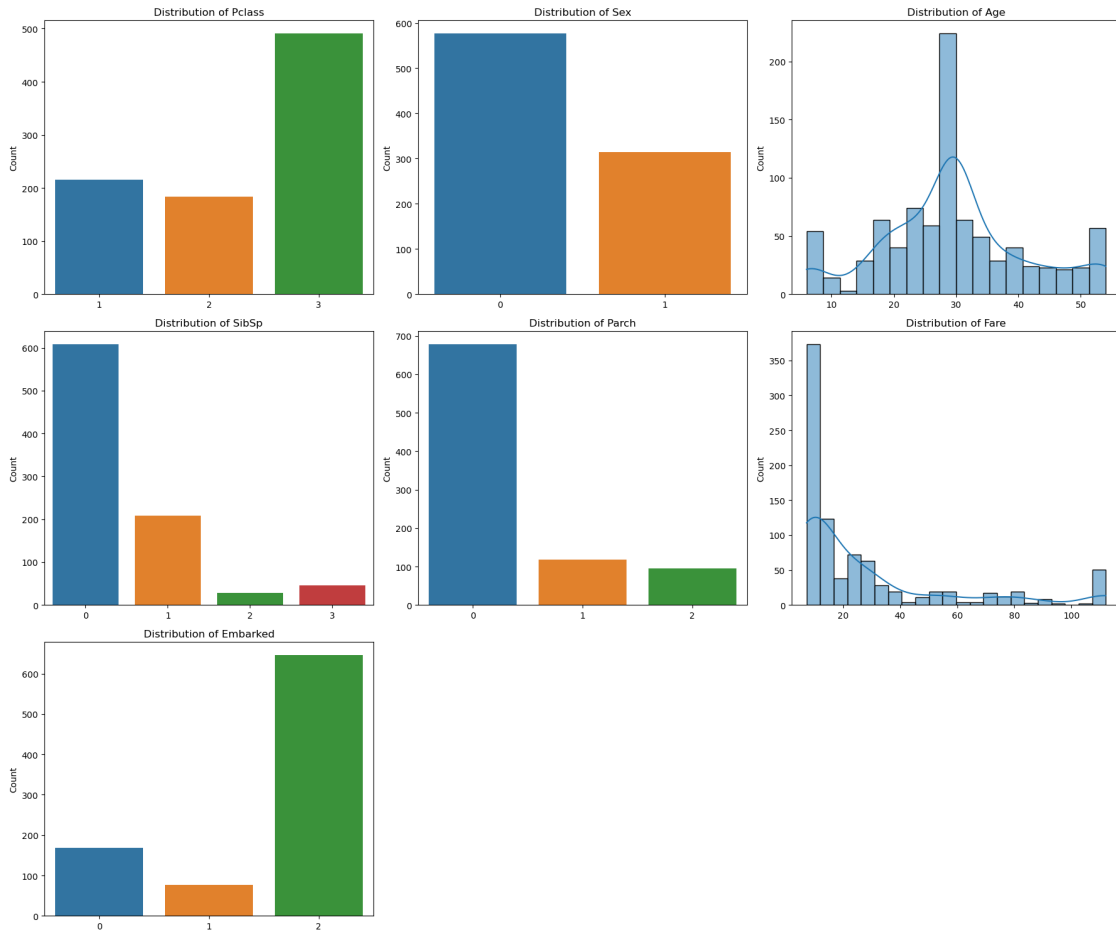
fig, axes = plt.subplots(nrows=3, ncols=3, figsize=(18, 15))

axes = axes.flatten()


for i, feature in enumerate(features):
    if df1[feature].dtype == 'object' or len(df1[feature].unique()) < 10:
        sns.countplot(x=feature, data=df1, ax=axes[i])
    else:
        sns.histplot(df1[feature], ax=axes[i], kde=True)
    axes[i].set_title(f'Distribution of {feature}')
    axes[i].set_xlabel('')
    axes[i].set_ylabel('Count')

for j in range(len(features), len(axes)):
    fig.delaxes(axes[j])


plt.tight_layout()
plt.show()
```



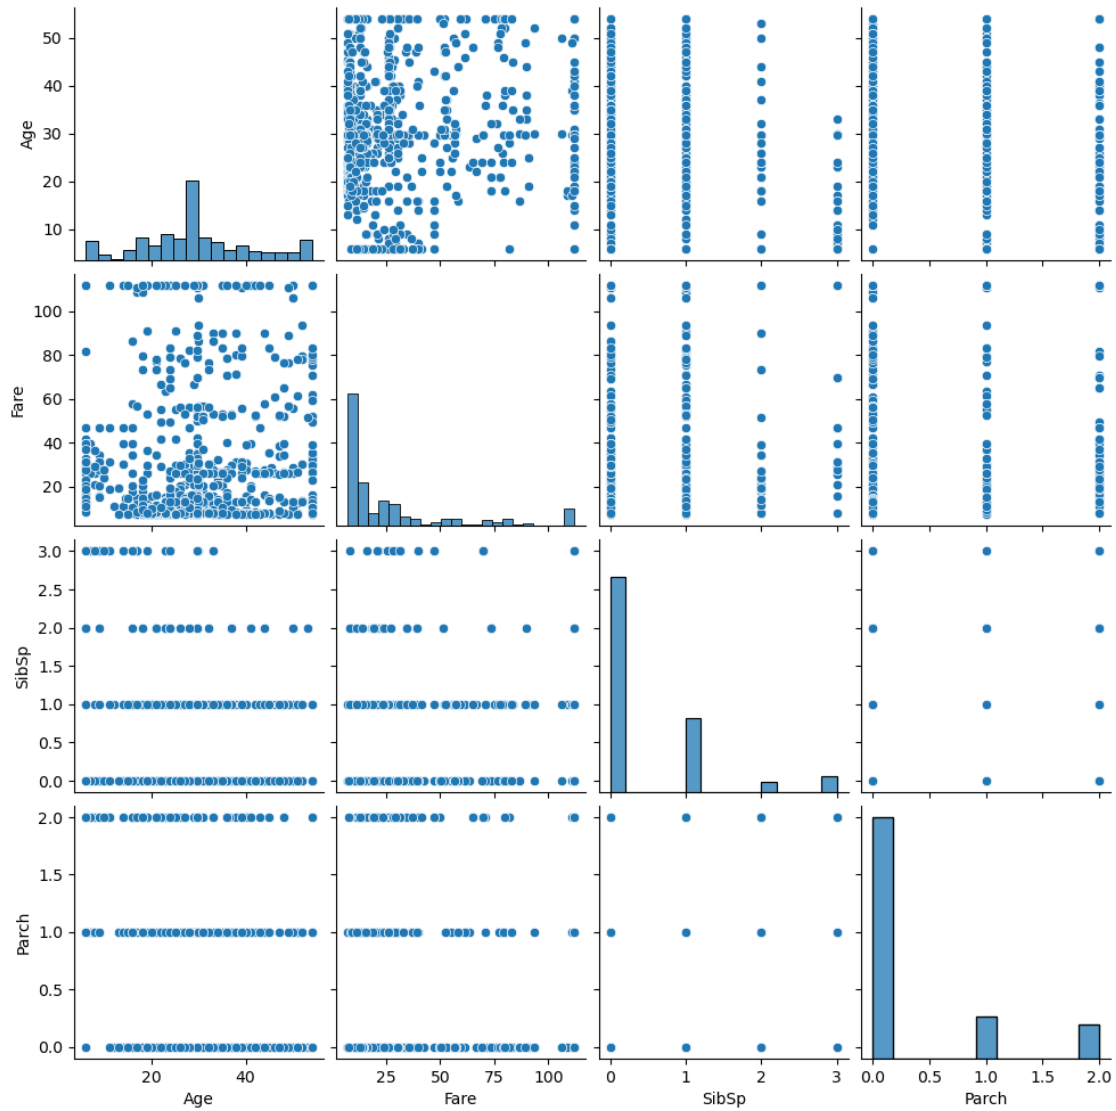
```
[39]: plt.figure(figsize=(10, 6))
sns.boxplot(x='Pclass', y='Age', data=df1)
plt.title('Age Distribution by Passenger Class')
plt.show()
```




```
[40]: plt.figure(figsize=(10, 6))
sns.boxplot(x='Pclass', y='Fare', data=df1)
plt.title('Fare Distribution by Passenger Class')
plt.show()
```




```
[41]: sns.pairplot(df1[['Age', 'Fare', 'SibSp', 'Parch']])  
plt.show()
```




```
[42]: plt.figure(figsize=(10, 6))
corr_matrix = df1.corr()
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', center=0)
plt.title('Correlation Matrix Heatmap')
plt.show()
```




```
[43]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Fare', hue='Survived', data=df1, palette='coolwarm')
plt.title('Scatter Plot of Age vs. Fare Colored by Survived')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.show()
```


```
[44]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Fare', data=df1)
plt.title('Scatter Plot of Age vs. Fare')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.show()
```



```
[45]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Fare', hue='Survived', data=df1, palette='coolwarm')
plt.title('Scatter Plot of Age vs. Fare Colored by Survived')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.show()
```



```
[46]: plt.figure(figsize=(10, 6))
sns.scatterplot(x='Age', y='Fare', hue='Pclass', data=df1, palette='viridis')
plt.title('Scatter Plot of Age vs. Fare Colored by Pclass')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.show()
```



```
[47]: fig, axs = plt.subplots(3, 2, figsize=(15, 18))

sns.regplot(x='Age', y='Fare', data=df, ax=axs[0, 0])
axs[0, 0].set_title('Linear Regression Plot of Age vs. Fare')

sns.scatterplot(x='Age', y='Fare', hue='Survived', data=df, palette='coolwarm',
               ↪ax=axs[0, 1])
sns.regplot(x='Age', y='Fare', data=df, scatter=False, ax=axs[0, 1])
axs[0, 1].set_title('Linear Regression Plot of Age vs. Fare with Hue by
               ↪Survived')

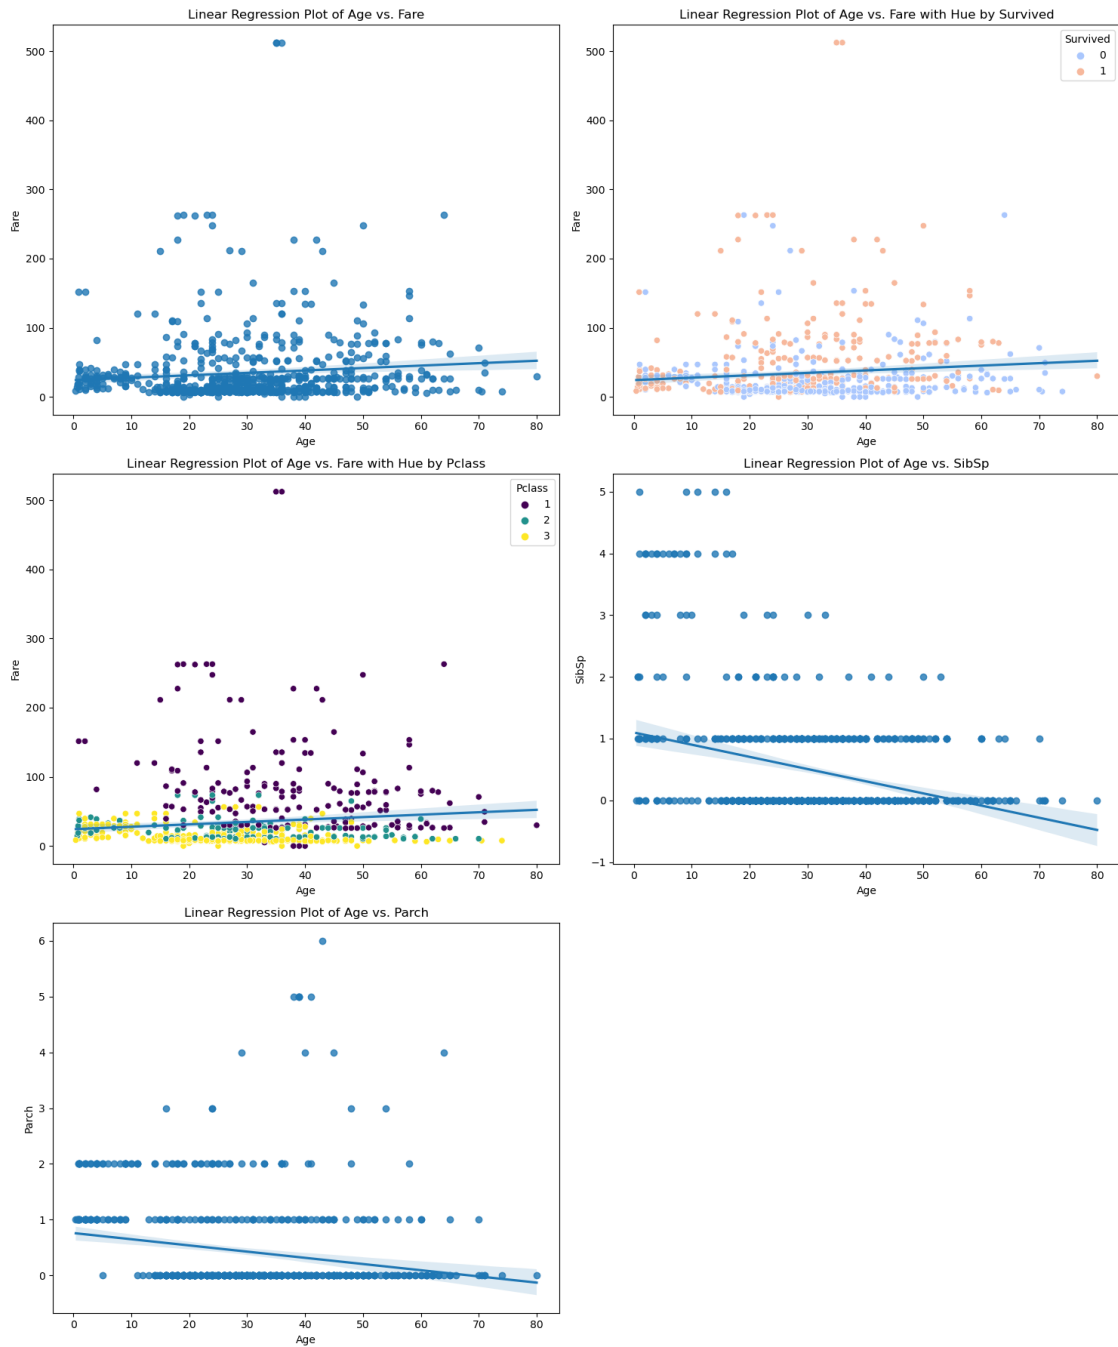
sns.scatterplot(x='Age', y='Fare', hue='Pclass', data=df, palette='viridis',
               ↪ax=axs[1, 0])
sns.regplot(x='Age', y='Fare', data=df, scatter=False, ax=axs[1, 0])
axs[1, 0].set_title('Linear Regression Plot of Age vs. Fare with Hue by Pclass')

sns.regplot(x='Age', y='SibSp', data=df, ax=axs[1, 1])
axs[1, 1].set_title('Linear Regression Plot of Age vs. SibSp')

sns.regplot(x='Age', y='Parch', data=df, ax=axs[2, 0])
axs[2, 0].set_title('Linear Regression Plot of Age vs. Parch')

axs[2, 1].axis('off')
```

```
plt.tight_layout()
plt.show()
```



```
[ ]:
```

```

[48]: x = df1.drop(['Survived'], axis=1)
      y = df1['Survived']

[49]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.2,
      ↪random_state=42)

[50]: x_train.shape

[50]: (712, 8)

[51]: x_test.shape

[51]: (179, 8)

[52]: scaler = StandardScaler()

[53]: x_train_scaled = scaler.fit_transform(x_train)
      x_test_scaled = scaler.transform(x_test)

[54]: x_train_scaled_df = pd.DataFrame(x_train_scaled, columns=x_train.columns)
      x_test_scaled_df = pd.DataFrame(x_test_scaled, columns=x_test.columns)

[55]: train_scaled = pd.concat([x_train_scaled_df, y_train.reset_index(drop=True)],
      ↪axis=1)
      test_scaled = pd.concat([x_test_scaled_df, y_test.reset_index(drop=True)],
      ↪axis=1)

[56]: from sklearn import preprocessing
      import matplotlib
      import matplotlib.pyplot as plt

      %matplotlib inline
      matplotlib.style.use('fivethirtyeight')
      # data
      x = pd.DataFrame({
          # Distribution with lower outliers
          'x1': np.concatenate([np.random.normal(20, 2, 1000), np.random.normal(1, 2,
      ↪25)]),
          # Distribution with higher outliers
          'x2': np.concatenate([np.random.normal(30, 2, 1000), np.random.normal(50,
      ↪2, 25)]),
      })
      np.random.normal

      scaler = preprocessing.RobustScaler()
      robust_df = scaler.fit_transform(x)
      robust_df = pd.DataFrame(robust_df, columns=['x1', 'x2'])

```

```

scaler = preprocessing.StandardScaler()
standard_df = scaler.fit_transform(x)
standard_df = pd.DataFrame(standard_df, columns=['x1', 'x2'])

scaler = preprocessing.MinMaxScaler()
minmax_df = scaler.fit_transform(x)
minmax_df = pd.DataFrame(minmax_df, columns=['x1', 'x2'])

fig, (ax1, ax2, ax3, ax4) = plt.subplots(ncols = 4, figsize =(20, 5))
ax1.set_title('Before Scaling')


sns.kdeplot(x['x1'], ax = ax1, color='r')
sns.kdeplot(x['x2'], ax = ax1, color='b')
ax2.set_title('After Robust Scaling')

sns.kdeplot(robust_df['x1'], ax = ax2, color='red')
sns.kdeplot(robust_df['x2'], ax = ax2, color='blue')
ax3.set_title('After Standard Scaling')

sns.kdeplot(standard_df['x1'], ax = ax3, color='black')
sns.kdeplot(standard_df['x2'], ax = ax3, color='g')
ax4.set_title('After Min-Max Scaling')

sns.kdeplot(minmax_df['x1'], ax = ax4, color='black')
sns.kdeplot(minmax_df['x2'], ax = ax4, color='g')
plt.show()

```



```

[57]: from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error , r2_score
lr = LinearRegression()
lr.fit(x_train,y_train)
y_predict = lr.predict(x_test)
R2score = r2_score( y_test , y_predict )*100
print('The R2score for Linear Regression is ' , R2score)

```

The R2score for Linear Regression is 43.496043988962654

```
[58]: from sklearn.tree import DecisionTreeRegressor
      # Instantiate model with 1000 decision trees

      dt = DecisionTreeRegressor()
      # Train the model on training data
      dt.fit(x_train, y_train);
      y_predict = dt.predict(x_test)
      print('The R2score for Decision Tree Regressor is ' , r2_score( y_test , y_predict ))
```

The R2score for Decision Tree Regressor is -0.15186615186615193

```
[59]: from sklearn.ensemble import RandomForestRegressor
      # Instantiate model with 1000 decision trees

      rf = RandomForestRegressor(n_estimators = 1000, random_state = 42)
      # Train the model on training data
      rf.fit(x_train, y_train);
      y_predict = rf.predict(x_test)
      print('The R2score for Random Forest Regressor is ' , r2_score( y_test , y_predict ))
```

The R2score for Random Forest Regressor is 0.4252145743886744

```
[60]: featureImportance = pd.Series(rf.feature_importances_ , index = x_train.
      ↪columns).sort_values(ascending = True)
      featureImportance
```

```
[60]: Parch          0.012892
      Embarked       0.024797
      SibSp          0.042184
      Pclass         0.101319
      Age            0.148472
      Fare           0.168210
      PassengerId     0.208006
      Sex            0.294120
      dtype: float64
```

```
[61]: df_x_train = x_train[ ['Pclass', 'Age', 'Fare']]
      df_x_test = x_test[ ['Pclass', 'Age', 'Fare']]
      rf_improved = RandomForestRegressor(n_estimators = 1000, random_state = 42)
      # Train the model on training data
      rf_improved.fit(df_x_train, y_train);
      y_predict = rf_improved.predict(df_x_test)

      R2score = r2_score( y_test , y_predict )
```



```
print('The R2score for Random Forest Regressor is ' , (R2score*100))
```

The R2score for Random Forest Regressor is 15.083418475481148

```
[62]: from sklearn.tree import DecisionTreeRegressor, DecisionTreeClassifier
dt = DecisionTreeRegressor(random_state = 42)
# Train the model on training data
dt.fit(x_train, y_train);
y_predict = dt.predict(x_test)

R2score = r2_score( y_test , y_predict )
print('The R2score for Decision tree Regressor is ' , (R2score *100))
```

The R2score for Decision tree Regressor is -5.9716859716859805

```
[63]: from sklearn.svm import SVR
svr_rbf = SVR(kernel='rbf', C=1e3, gamma=0.1)

y_rbf = svr_rbf.fit(x_train, y_train).predict(x_test)

R2score_rbf = r2_score( y_test , y_rbf )
print('The R2score for RBF is ' , R2score_rbf )
```

The R2score for RBF is 0.046543535936969915

```
[64]: from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, confusion_matrix, \
    classification_report
```

```
[65]: logreg = LogisticRegression(max_iter=1000)
logreg.fit(x_train, y_train)
```

```
[65]: LogisticRegression(max_iter=1000)
```

```
[66]: y_predict = logreg.predict(x_test)
accuracy = accuracy_score(y_test, y_predict)
print('Accuracy of Logistic Regression model:', accuracy * 100)
```

Accuracy of Logistic Regression model: 79.88826815642457

```
[67]: conf_matrix = confusion_matrix(y_test, y_predict)
print('Confusion Matrix:')
print(conf_matrix)
```

Confusion Matrix:

```
[[88 17]
 [19 55]]
```

```
[68]: class_report = classification_report(y_test, y_predict)
print('Classification Report:')
print(class_report)
```

Classification Report:

	precision	recall	f1-score	support
0	0.82	0.84	0.83	105
1	0.76	0.74	0.75	74
accuracy			0.80	179
macro avg	0.79	0.79	0.79	179
weighted avg	0.80	0.80	0.80	179

```
[69]: from sklearn.metrics import roc_auc_score, roc_curve, auc
logreg.fit(x_train, y_train)
```

```
[69]: LogisticRegression(max_iter=1000)
```


```
[70]: y_pred_proba = logreg.predict_proba(x_test)[:, 1]
```

```
[71]: roc_auc = roc_auc_score(y_test, y_pred_proba)
print('ROC AUC score:', roc_auc)
```

ROC AUC score: 0.8764478764478765

```
[72]: fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
```

```
[73]: plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
    roc_auc)
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



[74]:

[]:

[]: